


Lecture 1: Introduction

AC295

Advanced Practical Data Science
Pavlos Protopapas



Outline

1 : Why you should take this class and why not

2: Who are we

3: Course structure and activities

4: Expectations

5: Workload

6: Logistics

7: Grades

Why you should take this class

Because you want to learn how to:

- Put your model in production
- Integrate and orchestrate applications
- Deploy increasing amount of data
- Take advantage of available models
- Evaluate and debug model using visualization

If you have attended **ComputeFest** and found the topics interesting this class will also be interesting

Why you shouldn't take this class

You are **not** familiar with most of the concepts covered in CS109A/B

For example:

- Basic Machine Learning
- CNNs, RNNs, Autoencoders, GANs, etc
- Basic linux command

Remember, this course will be offered again in the fall!

UNK: from Data Science Series to Real World

Data Science Series

Real World

Ask Question

Collect Data

CSV file, images,
scraping



Manage larger database

EDA

Notebook



Learn packages to process larger amount of data

Methodology

Multiple tasks




Handle complex team dynamics and orchestrate applications


Story-telling

Webpage, blogs,
posts

UNK: Real World Example



UNK: Real World Example EDA



Who?

Pavlos Protopapas

Teaches CS109(a/b), the data science capstone course, and AC295 (advanced practical data science). Research in astrostatistics: machine learning, statistical learning, big data for astronomical problems.

He has picked some new hobbies besides 109 and **eating**:

Going to BSO (see you there), cross country ski (completed Engadin skimarathon), cheese making and being a TikToker (check me out @pavlosprotopapas)



Who? (cont)

Michael S. Emanuel

After 17 years in finance, mainly fixed income portfolio management, Michael started a second career and is completing the Masters of Data Science program at Harvard. He is a father of two small children who occasionally crash IACS events and enjoys distance running and classical music.



Who? (cont)

Andrea Porelli

Urban planner turned into data hacker. He likes to break things just for the sake of putting them back together (most of the time). Committed to apply Data Science to change something. So far, he managed to change himself the most -thanks IACS- and look forward to pass it over.



Who? (cont)

Giulia Zerbini

Data Designer. Creative technologist at *The Visual Agency* in Milan, MA Graduate at Politecnico di Milano. Designing and developing visualizations and interfaces based on data. Passionate about using visualizations for discovering patterns in data and communicating information in intuitive terms to a broad audience.



Course Structure and Activities

Modules:

1. Deploy data science (integration + scalability)
2. Transfer learning and distillation
3. Visualization as investigative tool

Activities:

lectures, reading discussions, exercises, practicums, projects

Lectures: Tuesday and Thursday 4:30-5:45 pm in Cruft 309

Office Hours: TBD

Topics

Deploy data science (integration + scalability)

- A. Virtual Environments, Virtual Boxes, and Containers
- B. Kubernetes
- C. Dask

Topics (cont)

Transfer learning and distillation

- A. Intro to Transfer Learning: Basic Transfer Learning and SOTA Models
- B. Transfer Learning across Tasks
- C. Distillation and Compression

Topics (cont)

Visualization as investigative tool

- A. Introduction and Overview of Viz for Deep Models
- B. Convolutional Neural Networks for Image Data
- C. Recurrent Neural Networks for Text Data

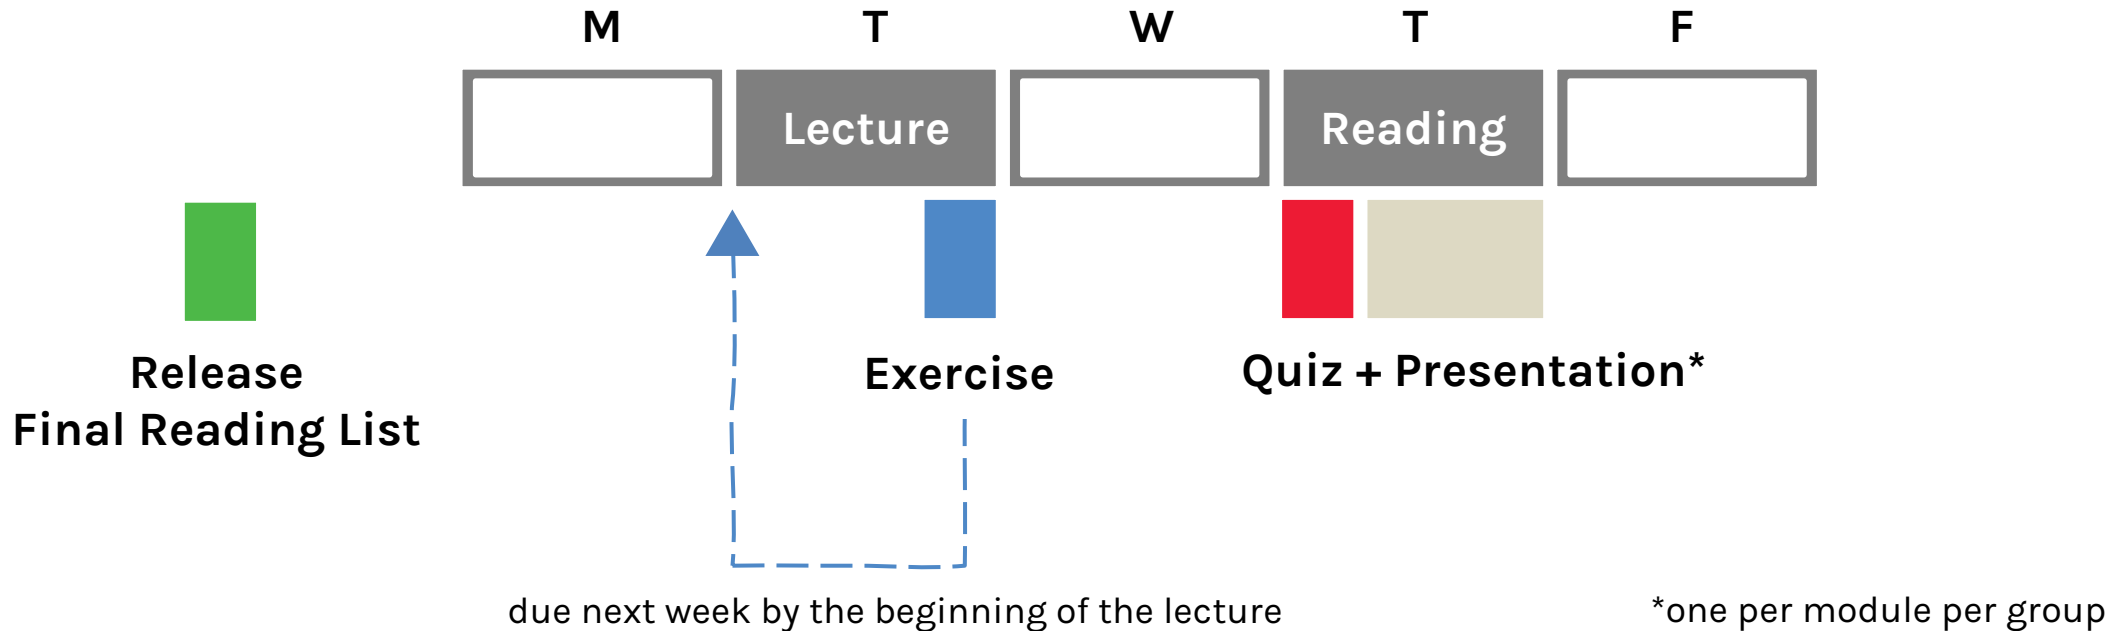
Calendar

> [Link to Calendar](#) <

Week	Date	Lecture #	Topics
1	1/28	1	Introduction
	1/30	2	Set up
2	2/4	3	Virtual Environments, Virtual Boxes, and Containers
	2/6	4	Journal Discussion
3	2/11	5	Kubernetes
	2/13	6	Journal Discussion
4	2/18	7	Dask
	2/20	8	Journal Discussion
5	2/25	9	Practicum
	2/27	10	Practicum
6	3/3	11	Intro to Transfer Learning: Basic Transfer Learning and SOTA Models
	3/5	12	Journal Discussion
7	3/10	13	Transfer Learning across Tasks
	3/12	14	Journal Discussion
	3/17	Spring Recess	
	3/19		
8	3/24	15	Distillation and Compression
	3/26	16	Journal Discussion
9	3/31	17	Practicum
	4/2	18	Practicum
10	4/7	19	Introduction and Overview of Viz for Deep Models
	4/9	20	Journal Discussion
11	4/14	21	Convolutional Neural Networks for Image Data
	4/16	22	Journal Discussion
12	4/21	23	Recurrent Neural Networks for Text Data
	4/23	24	Journal Discussion
13	4/28	25	Practicum
	4/30	26	Practicum
14	5/5	27	Project
	5/7	28	Project
15	5/12	29	Project
	5/14	30	Project

Course Structure and Activities

Regular week schedule



Workload

Regular Week

3 hours in class
3 hours reading
2 hours exercise
4 hours presentation*

~ 12 hours/week

* 1 presentation per module per group (3 total)

Practicum and Project Week

~ 15 hours/week**

** 3 practicums and 1 final project (2 weeks long)

We will be asking for your feedback on the workload

Expectations

How to read and present class material

> [Link to Reading Guidelines](#) <

> [Link to Presentation Guidelines](#) <

Logistics

Fill up forms

Make group*

Sign-up presentation**

* Fill group components in each row

** Each group should pick one slot (white background) in each module

7: Grades

Assignment	Final Grade Weight
Quizzes	9%
Exercises	9%
Presentations	15%
Practicums	45%
Projects	20%
Participation	2%
Total	100%

This is the first time we are giving the course, so your feedback will be vital in tuning it this year and improving it for future years.

However, we are making every effort to have a well organized course with as little glitches as possible and we promise you an exciting semester

AC295

Advanced Practical Data Science

Pavlos Protopapas